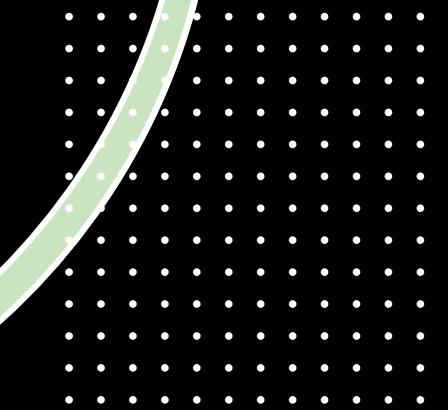


Data cleaning and
transformation



Agenda

- Handle Missing Data & Duplicates
- Merge, Join, Concatenate DataFrames
- Apply Aggregation & GroupBy
- Generate Summary Stats & Insights





What is "Messy" Data

- Real-world data is almost never clean. It often has:
- **Missing Data:** Represented as NaN (Not a Number). This can be from data entry errors, or data that was never collected.
- **Duplicates:** The same row appearing more than once.

Why does it matter?

Missing data can break your calculations (e.g., `mean()`, `sum()`).

Duplicates can skew your results (e.g., making sales look higher than they are).

Sample data

- data = {'A': [1, 2, np.nan, 4],
 'B': [5, np.nan, 7, 8],
 'C': [9, 10, 11, 12]}

```
df = pd.DataFrame(data)  
print(f"Original Data:\n{df}\n")
```

df.isnull()

Find missing data
(returns
True/False)



count of missing
data per column

`df.isnull().sum()`



Drop rows with
any NaN

`df.dropna()`

Fill NaN
with a value

`df.fillna(value=0)`

Fill with the column's mean

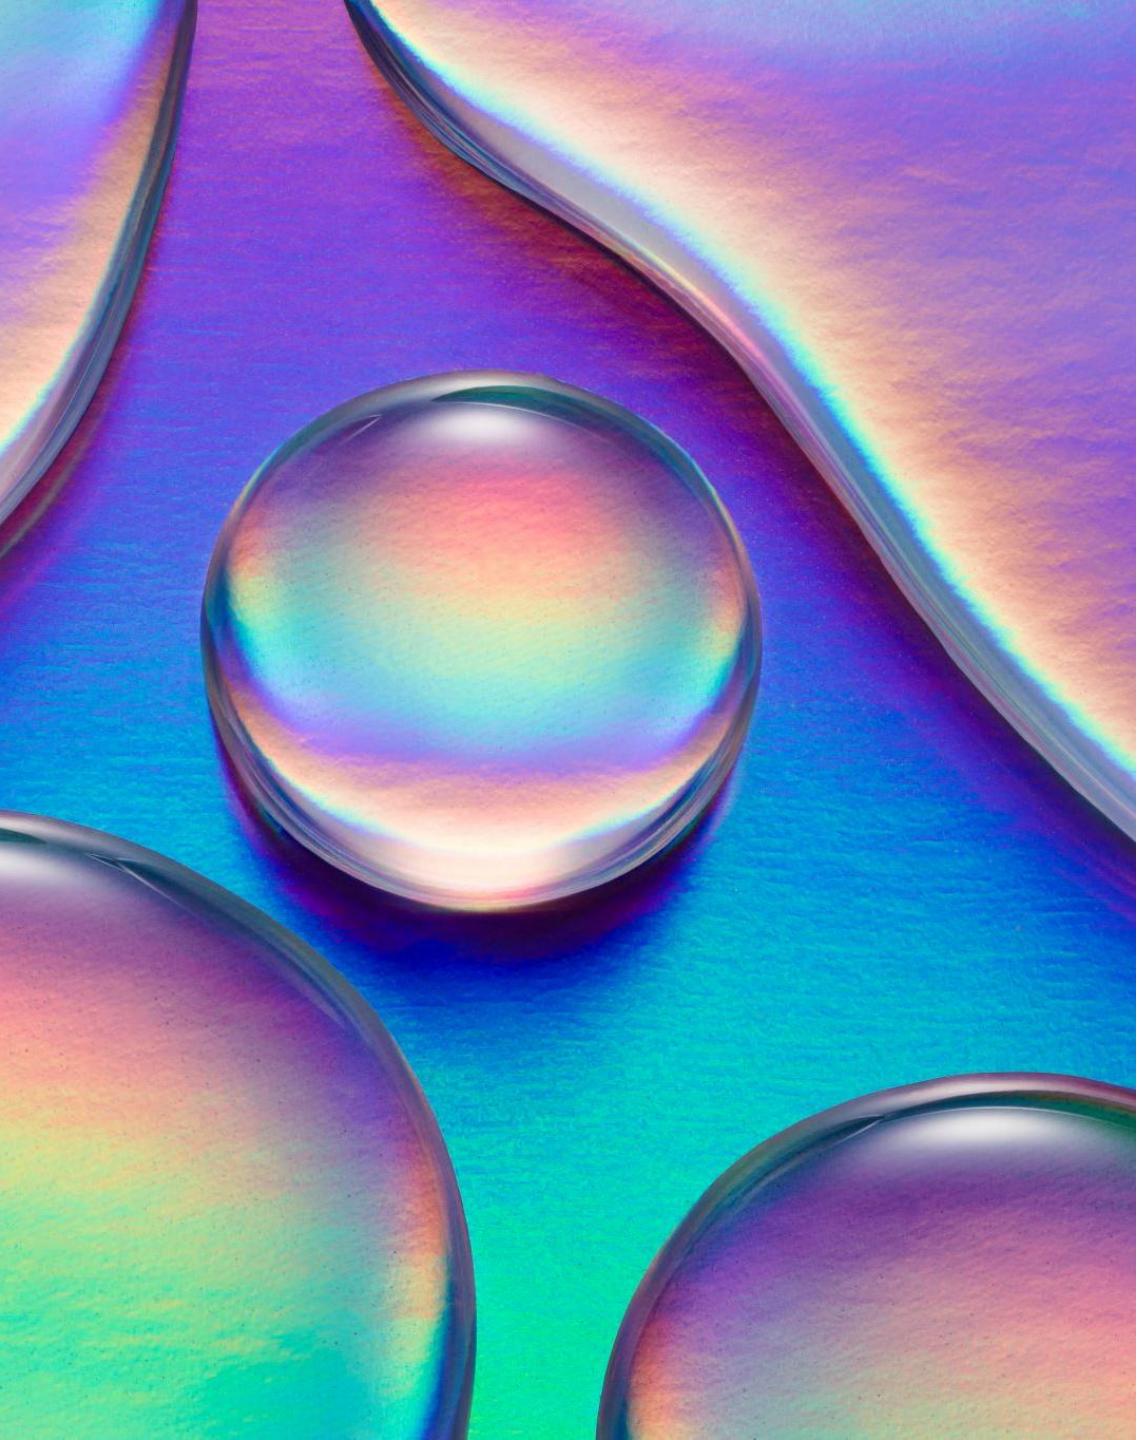
- `mean_b = df['B'].mean()`
- `df['B'] = df['B'].fillna(value=mean_b)`

Finding & Handling Duplicates

- `DataFrame.duplicated(subset=None, keep='first')`
- Return boolean Series denoting duplicate rows.

```
duplicated(subset=None,  
          keep='first')
```

- **Subset column label or sequence of labels, optional**
- Only consider certain columns for identifying duplicates, by default use all of the columns.
- **Keep {‘first’, ‘last’, False}, default ‘first’**
- Determines which duplicates (if any) to mark.
- first : Mark duplicates as True except for the first occurrence.
- last : Mark duplicates as True except for the last occurrence.
- False : Mark all duplicates as True.



Drop duplicate rows

- **DataFrame.drop_duplicates(subset=None, *, keep='first', inplace=False, ignore_index=False)**
- df.drop_duplicates()
-

`.drop_duplicates()`

- **`inplace`bool, default False**
 - Whether to modify the DataFrame rather than creating a new one.
- **`ignore_index`bool, default False**
 - If True, the resulting axis will be labeled 0, 1, ..., n - 1.

Hands-On

```
raw_data = {  
    'Employee': ['John', 'Anna', 'Peter',  
    'John', 'Linda', 'Anna'],  
    'Department': ['Sales', 'IT', 'Sales',  
    'Sales', 'HR', 'IT'],  
    'Salary': [50000, 75000, 52000,  
    50000, 60000, np.nan]  
}  
  
emp_df = pd.DataFrame(raw_data)
```

- Task 1:Find and print the total count of missing values for each column.
- Task 2:Drop the duplicate rows (Hint: 'John' and 'Anna' appear twice).
- Task 3 (Bonus):Fill the missing Salary with the average salary of the other employees.

Aggregation and GroupBy



“Split-Apply-Combine” Strategy

- Split: You groupby color (whites, darks, colors).
- Apply: You apply a function to each group
- Combine: You combine the results



Split

Split: df.groupby('Department') -> splits data into groups (Sales, IT, HR).

Apply

`['Salary'].mean()` -> applies the "mean" function to the
'Salary' of each group.

Example

- data = {'Employee': ['John', 'Anna', 'Peter', 'Linda'], 'Department': ['Sales', 'IT', 'Sales', 'HR'], 'Salary': [50000, 75000, 52000, 60000]}
- df = pd.DataFrame(data)
- # 1. Find the AVERAGE salary for each department
- # 2. Find the SUM of salaries for each department
- # 3. Find the COUNT of employees in each department
- # 4. Using .agg() for multiple operations

aggregate

- it's used to apply one or more aggregation functions to your grouped data.
- Apply **multiple functions** at once (e.g., mean, max, count)
- Rename the output columns

Hands-On Code: Groupby

- sales_data = {
 - 'Region': ['North', 'South', 'North', 'South', 'East', 'East'],
 - 'Product': ['A', 'A', 'B', 'B', 'A', 'B'],
 - 'Amount': [100, 150, 200, 250, 50, 300]
- }
- df_sales = pd.DataFrame(sales_data)
- 1. What is the *total* sales Amount for each Region?
- 2. What is the *average* sales Amount for each Product?
- 3. What is the *total* sales Amount for each *combination* of Region and Product?
(Hint: You can groupby a list: ['Region', 'Product'])

Getting Quick Insights

- Now that your data is clean and organized, you can get high-level insights.

`.describe()` - Stats for all numeric columns

- Gives count, mean, std dev, min, max, quartiles

`.value_counts()`

Frequency of unique values #
Great for categorical data

.info()

.corr() - Correlation

- Shows how strongly numeric columns are related
- Need more numeric columns for this to be useful
- `df.corr()`