


Pandas

- The Pandas library is built on NumPy and provides easy-to-use data structures and data analysis tools for the Python programming language

Installing pandas

- pip install pandas
- %pip install pandas

- import pandas as pd
- print(pd.__version__)



Pandas Data Structures

- Series (1D labeled array)
- DataFrame (2D table with rows and columns)

Pandas series

A one-dimensional
labeled array capable
of holding any data type

```
s = pd.Series([3, -5, 7,  
4], index=['a', 'b', 'c', 'd'  
])
```

question

- Tracking Monthly Sales
- sales = pd.Series([5000, 7000, 8000, 6500, 7200], index=['Jan', 'Feb', 'Mar', 'Apr', 'May'])
- Find February month sale
- Find the average sale
- Days sale above 6000
- Day with the highest sales
- New Series with 10% increase in sales

Dataframe

- A two-dimensional labelled data structure with columns of potentially different types
- `data = {`
- `'Name': ['Alice', 'Bob', 'Charlie', 'David'],`
- `'Age': [25, 30, 35, 40],`
- `'Salary': [50000, 60000, 70000, 80000],`
- `'Department': ['HR', 'IT', 'IT', 'Finance']`
- `}`
- `df = pd.DataFrame(data)`

Manipulate: Add column

- import pandas as pd
- # Mock data: Coffee shop customers
- data = {
- 'Name': ['Alice', 'Bob', 'Charlie'],
- 'Age': [25, 30, 35],
- 'Spend': [45.0, 30.0, 50.0]
- }
- df['Loyalty'] = df['Spend'] > 40 # Boolean
- print("\nUpdated:\n", df)

Reading and Writing Files

- Loading Sales Data from CSV
- `df = pd.read_csv('sales.csv')`
- `print(df.head()) # Display first 5 rows`

Writing csv

- df.to_csv('new_sales.csv', index=False)

Reading file in Numpy

- With no missing values
- Use **numpy.loadtxt**
- With missing values
- Use **numpy.genfromtxt**.

Pandas Indexing:

- `.loc`
 - `.loc` (label-based)
 - Row/column names (labels)
 - `df.loc[row_label, col_label]`

iloc

- .iloc (position-based)
- Row/column integer positions
- df.iloc[row_pos, col_pos]

Sample data

- import pandas as pd
- data = {
 - 'Customer': ['Alice', 'Bob', 'Charlie', 'Diana'],
 - 'Drink' : ['Latte', 'Espresso', 'Cappuccino', 'Americano'],
 - 'Price' : [4.50, 3.00, 5.00, 2.50],
 - 'Qty' : [2, 3, 1, 4]
- }
- df = pd.DataFrame(data, index=['A1', 'A2', 'A3', 'A4'])

- 1. From csv take Single row(Drink) by label
- 2. Multiple rows from Drink to Price

.loc

- # 1. Single row by label
- df.loc['A2']
- # 2. Multiple rows (inclusive)
- df.loc['A1':'A3']
- # 3. Specific cells
- df.loc['A1', 'Price'] # → 4.50
- df.loc[['A1','A4'], ['Drink','Qty']]
- # 4. Boolean mask with labels
- df.loc[df['Price'] > 3.5, 'Customer']

iloc

- # 1. First 2 rows (positions 0,1)
 - df.iloc[:2]
- # 2. Row 1 to 3 (exclusive end)
 - df.iloc[1:4]
- # 3. 2nd & 4th rows, 1st & 3rd columns
 - df.iloc[[1,3], [0,2]]
- # 4. Last column only
 - df.iloc[:, -1]

dropping

- `s.drop(['jan']) #Drop values from rows`
- `Df.drop("name") #Drop values from columns(axis=1)`

Data Cleaning

- Handling Missing Data
- df = pd.DataFrame({
 - 'Customer': ['A', 'B', 'C', 'D'],
 - 'Age': [25, None, 35, None],
 - 'Purchase': [500, 1000, None, 700]
- })
- df.fillna(df.select_dtypes(include=['number']).mean(), inplace=True)
- # Fill missing values with column mean
- print(df)

Data Filtering

- Find Customers Who Spent More Than 800
- `high_spenders = df[df['Purchase'] > 800]`
- `print(high_spenders)`



select elements
from an array based
on multiple
conditions

- `np.where()`
- `np.where(condition, x, y)`
 - x: The value (or array) to use where the condition is True
 - y: The value (or array) to use where the condition is False



Solve Day 1

- test_RecallPandasDataFrame
- **test_WorkwithNumPy2DArray**

Convert Pandas DataFrame to NumPy Array

- Why convert?
Use NumPy for fast math, machine learning input, or plotting.
- `.to_numpy()`
- `df.to_numpy()`



function converts Series to DataFrame

.to_frame()

- Series.**to_frame()**

Solve

- Practice question 1 and 3
- test_UnderstandStructuredData1/2
- test_SliceandModify2DArrays

Grouping and Aggregation

- Grouping and aggregation allow us to summarize data based on certain categories.
- This is useful for reports, performance analysis, and trend identification.



```
• # Creating a sample employee dataset
• df = pd.DataFrame({
•     'Employee': ['Alice', 'Bob', 'Charlie', 'David', 'Eve',
•                  'Frank'],
•     'Department': ['HR', 'IT', 'IT', 'c', 'HR', 'Finance'],
•     'Salary': [50000, 70000, 80000, 60000, 52000, 65000]
• })
•
• # Average salary per department
• avg_salary = df.groupby('Department')['Salary'].mean()
• print(avg_salary)
•
```

More Aggregations

- Finding Total Salary per Department
- Finding Number of Employees per Department

Task

- df = pd.DataFrame({
 • 'City': ['New York', 'Los Angeles', 'Chicago', 'New York', 'Chicago', 'Los Angeles'],
 • 'Sales': [200, 300, 250, 400, 350, 500],
 • 'Profit': [20, 40, 35, 50, 45, 60]
 • })
- **Write code to calculate:**
 - 1.Total sales per city.
 - 2.Average profit per city.
 - 3.Number of transactions per city.

Merging and Joining in Pandas

- Merging allows you to combine multiple datasets based on a common column.

example

- A company wants to **combine employee details with performance ratings.**
- df1 = pd.DataFrame({
 - 'ID': [1, 2, 3, 4],
 - 'Name': ['Alice', 'Bob', 'Charlie', 'David'],
 - 'Department': ['HR', 'IT', 'IT', 'Finance']
- })

- df2 = pd.DataFrame({
 - 'ID': [1, 2, 3, 4],
 - 'Performance': ['A', 'B', 'A', 'C'],
 - 'Bonus': [1000, 500, 1200, 300]
- })

- merged_df = pd.merge(df1, df2, on='ID')
- print(merged_df)

Types of Joins

- **Inner Join (Default)**
- Only keeps matching rows.
- `pd.merge(df1, df2, on='ID', how='inner')`

Types of Joins

- Left Join (Keep all from Left)
- pd.merge(df1, df2, on='ID', how='left')

Types of Joins


- Right join (keeps all from right)
- pd.merge(df1, df2, on='ID', how='left')

Types of Joins

- Outer join (keeps everything)
- `pd.merge(df1, df2, on='ID', how=outer')`


Merge, Join, and Concatenate

how='inner'




natural join

how='outer'




full outer join

how='left'



left outer join

how='right'



right outer join

Merge, Join, and Concatenate

Method	Use When
<code>pd.concat()</code>	Stack vertically/horizontally
<code>df.merge()</code>	SQL-like join on columns
<code>df.join()</code>	Join on index

Stacking Rows

- df1 = pd.DataFrame({'A': ['A0', 'A1'], 'B': ['B0', 'B1']})
- df2 = pd.DataFrame({'A': ['A2', 'A3'], 'B': ['B2', 'B3']})
- # 1. Stacking Rows (axis=0 is default)
- row_concat = pd.concat([df1, df2])
- row_concat_fixed = pd.concat([df1, df2], ignore_index=True)

Stacking Columns (axis=1)

- df3 = pd.DataFrame({'C': ['C0', 'C1'], 'D': ['D0', 'D1']})
- col_concat = pd.concat([df1, df3], axis=1)
- print(f"Column concat (axis=1):\n{col_concat}")

df.join()

- Join columns with *other* DataFrame either on index or on a key column.
- Efficiently join multiple DataFrame objects by index at once by passing a list

- >>> df = pd.DataFrame({'key': ['K0', 'K1', 'K2', 'K3', 'K4', 'K5'],
• ... 'A': ['A0', 'A1', 'A2', 'A3', 'A4', 'A5']})
- >>> other = pd.DataFrame({'key': ['K0', 'K1', 'K2'],
• ... 'B': ['B0', 'B1', 'B2']})
- df.join(other, lsuffix='_caller', rsuffix='_other')

Hands-On Code

- # List of orders
- df_orders = pd.DataFrame({
- 'order_id': [101, 102, 103, 104],
- 'customer_id': [1, 2, 1, 3]
- })

- # List of customers
- df_customers = pd.DataFrame({
- 'customer_id': [1, 2, 3, 4],
- 'name': ['Alice', 'Bob', 'Charlie', 'David']
- })

Task 1: Merge df_orders (left) with df_customers (right) to get the name for each order_id.(Hint: The key is customer_id. Use a left join).

Task 2 (Bonus): Which customer placed no orders?(Hint: Use how='outer' and look for NaNs).

Task

- products = pd.DataFrame({
 - 'Product_ID': [101, 102, 103, 104],
 - 'Product_Name': ['Laptop', 'Phone', 'Tablet', 'Monitor']
- })
- sales = pd.DataFrame({
 - 'Product_ID': [101, 103, 104, 105],
 - 'Units_Sold': [10, 5, 8, 3]
- })
- Perform an inner join to find which products were sold.
- Perform a left join to keep all products, even if they have no sales.
- Perform an outer join to keep all products and sales, even if they don't match.

